



Optimization Spark Application Using Scala

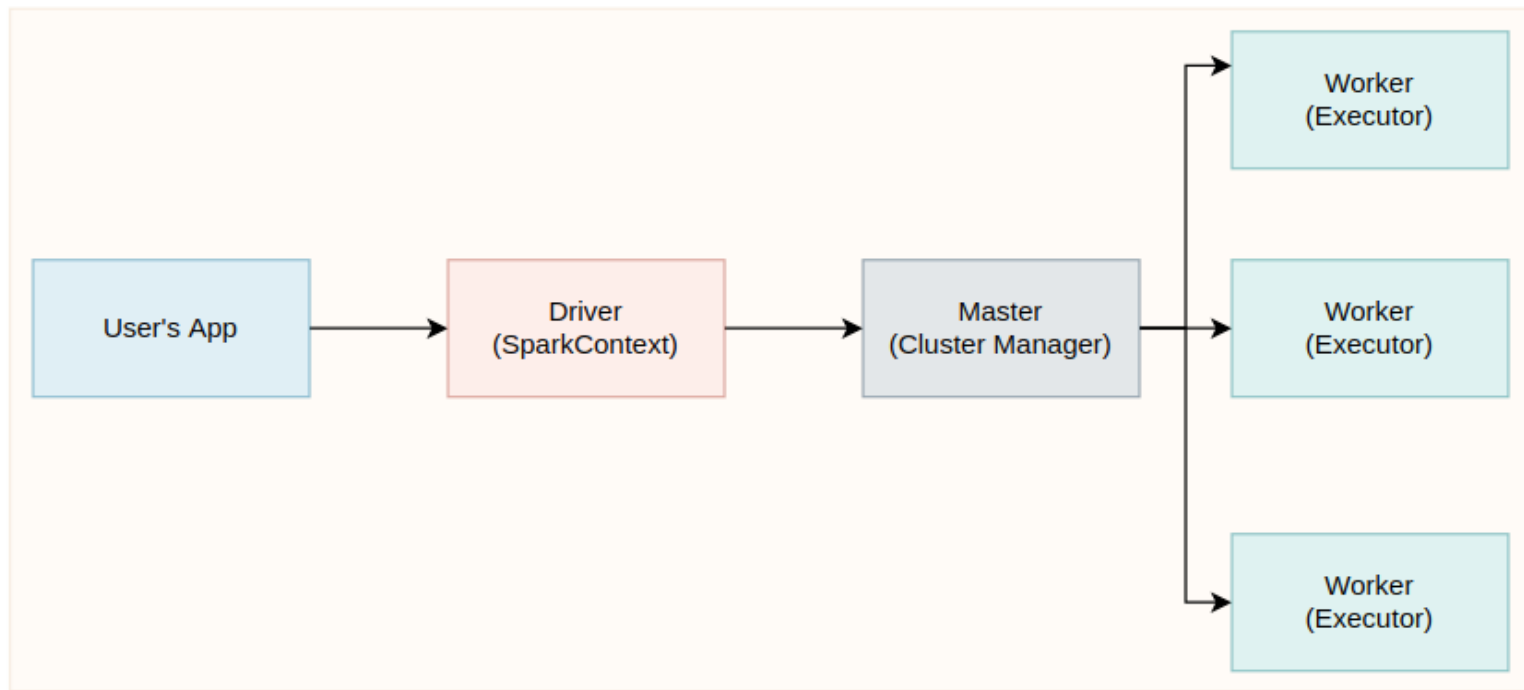
Kodjo Klouvi | kodjo@osekoo.com



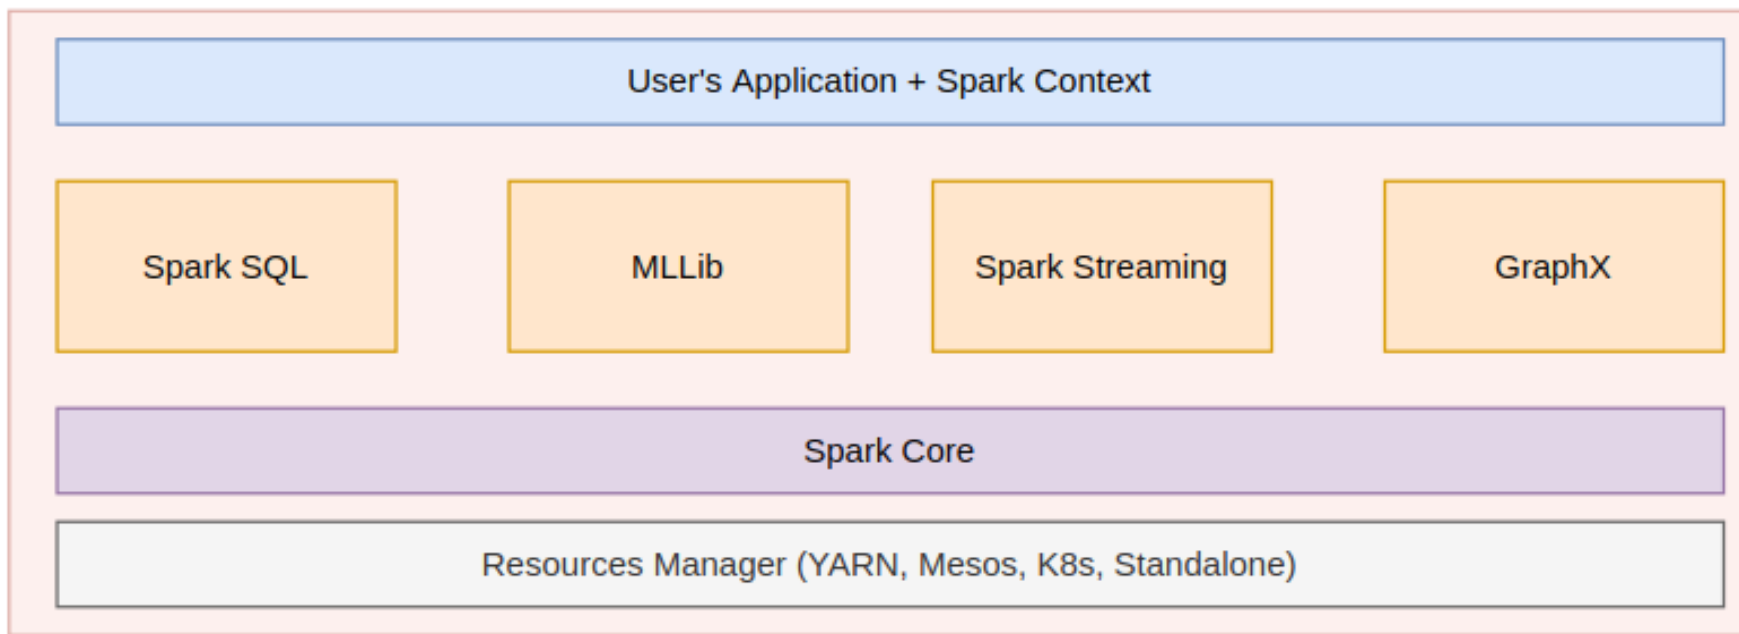
Apache Spark

- Big Data processing Engine (batch and streaming)
- In-memory data pipeline (ETL)
- Parallel execution
- Machine Learning

Spark Architecture



Spark Framework





Spark Optimization

- Make a difference between Transformations and Actions
- Use appropriate data format
- Clean data before Spark operations
- Filter data at the beginning of the ETL
- Avoid as much as possible Shuffle operations
- Apply appropriate partitioning



Spark Optimization

- Control the join tasks
- Use of broadcast join
- Cache data in memory or on disk
- Review logical and physical plans
- Be aware of DAGs
- Prefer Scala to Python
 - Python doesn't handle correctly UDFs
- Persist ML models' state



Transformations vs Actions

- Transformations = Lazy Evaluation
 - GroupBy, select, orderBy, join, filter, read
 - Only the logical plan is computed
 - Few computing resources are used at this stage
- Actions = Execution
 - Collect, show, count, save, take
 - Optimized and physical plans are computed
 - Transformations are computed
 - Very high consumption of resources



Data format

- Store/persist data in parquet format (column-oriented structure)
 - Suitable for big data tasks
 - For future tasks
 - After shuffle tasks (to avoid re-shuffle)
 - For machine learning tasks



Pre-cleaning Data

- Remove unnecessary data before start spark operations
- Drop columns
- Review the dataset and keep only useful data



Filtering data at the beginning

- Apply filter, groupByKey, reduceByKey, join as soon as possible
- Cache or persist the data in memory or on the disk depending on the data size



Shuffle or not...

- The most expensive action in Spark!
- Moving data over the cluster
- Use the appropriate strategy when filtering or joining dataset
 - Parallelize tasks (partitioning)
 - Broadcast
 - hash



Partitioning

- Partition data to parallelize tasks
 - Repartition, partitionBy, coalesce
- Partition data before saving it to disc



Review computation plans

- Spark computes 3 plans before execution
 - Logical plan
 - Optimized plan
 - Physical plans
- Explain(extend=true) displays all plans + DAGs
- Review the physical plan nad DAGs to detect any redundancy or unnecessary operations



Scala

- Object-oriented programming language
- Coupled with Java (Java bytecode/Gateway)
- Statically typed (String, Boolean, Int, Long, Float)
- SuperTypes **Any** and **AnyRef**



Scala

- Multithreading / Parallel run
- Supports optional parameters, named parameters, etc.
- Widely used by big companies (APple, Twitter, Google,...)
- Spark is built in Scala



Scala

- Case-sensitive: **DataFrame** is different from **Dataframe**
- Classes and interfaces name are camelcase with first capital letter (**HousePriceEstimator**)
- Methods and variables name are camelcase with first small letter (**getPrice()**)



Scala

- Use **val** keyword to create an **immutable** variable
 - `val lastUpdate = "2020-11-23"`
 - `val level : Int = 17`
- Use **var** keyword to create a **mutable** variable
 - `var gender = "F"`
 - `var strike: Float = 85.9`



Scala

- Example and comment at <https://github.com/osekoo/hands-on-spark-scala/blob/develop/get-started/src/main/scala/WordCount.scala>



Hands-On

- <https://github.com/osekoo/hands-on-spark-scala>
- Preparing dev environment
- Writing code
- Building/testing
- Packaging and submitting tasks





IntelliJ

- Integrated Development Environment
- Multiple modules available for Scala, SBT, Big Data, etc.
- Easy to debug
- Profiler available



Docker

- Platform as a service (PaaS)
- OS abstraction
- Self-contained applications running in containers
- Simplify applications delivery/distribution
- Enable (auto)scaling



Lab Session

- Sync Spark, Hadoop, Scala, Java and SBT versions
- We use here Spark 3.0.2, Hadoop 2.7, Scala 2.12.x, Java 8 and SBT 1.x
- Other details on the Github page



Troubleshooting

- Raise any question during the session or via email kodjo@osekoo.com